


Base Architecture Diagrams

Lucid Chart - Base Arch



Kolate System Architecture Documentation

Overview

This architecture diagram depicts a scalable, multi-tenant cloud-based system named "Kolate" (likely a custom application or platform), deployed on Amazon Web Services (AWS). The design follows a microservices-oriented pattern with emphasis on event-driven processing via Apache Kafka, separation of concerns between core services and tenant-specific data, and robust security/isolation through VPCs and load balancing.

The system supports high availability, fault tolerance, and horizontal scaling, utilizing AWS managed services for networking, compute, storage, and messaging. It appears to handle enterprise information provisioning, with dedicated infrastructure for shared core components and isolated tenant databases to ensure data sovereignty.

Key characteristics:

- **Deployment Model:** Multi-tenant SaaS (Software as a Service) with tenant-specific databases.
 - **Core Paradigm:** Microservices architecture with asynchronous messaging (Kafka).
 - **Scalability:** Leverages auto-scaling EC2 instances, load balancers, and containerized deployments.
 - **Security:** VPC isolation, API Gateway for ingress control.
 - **Data Management:** Hybrid of relational (PostgreSQL), NoSQL (MongoDB), and caching (Redis) databases, split between core and tenant layers.

The diagram is divided into logical zones: Client Ingress, Networking/Load Balancing, Application Layer (EC2 Apps), Core Microservices, Databases & Services, and Tenant-Specific Storage.

High-Level Data Flow

1. Client Request Ingress:

- External clients (e.g., web/mobile apps) resolve domains via Amazon Route 53.
- Traffic is cached and accelerated globally via Amazon CloudFront (CDN).
- Requests are routed to regional edge locations and balanced across instances using Elastic Load Balancing (ELB).

2. API Exposure:

- Requests hit the API Gateway, which handles authentication, throttling, caching, and routing to backend services.

3. Application Processing:

- API Gateway forwards to EC2 instances in the "Apps" cluster, hosting Kolate Predict Microservices.
- These microservices interact with the core Kafka-based event bus for orchestration.

4. Event-Driven Core:

- Kolate Core Microservices consume/produce events via Kafka (with Zookeeper for coordination).
- Core services query/provision data from shared Kolate Core Databases.

5. Enterprise Provisioning:

- The Enterprise Info Provider Service (in Databases & Services EC2) integrates with tenant data for provisioning workflows.
- Tenant-specific operations route to isolated EC2 Tenant Databases.

6. Data Persistence:

- Core data stored in shared databases (Redis for caching, MongoDB for documents, PostgreSQL for transactions).
- Tenant data isolated in per-tenant instances of the same DB stack, containerized for scalability.

7. Egress/Feedback:

- Responses flow back through the same ingress path, with potential async callbacks via Kafka.

This flow ensures loose coupling, resilience (e.g., via message queues), and compliance (e.g., tenant data isolation).

Components

The architecture is organized into layered components. Below is a detailed breakdown, grouped by zone.

1. Networking and Security Layer

Component	Description	Purpose	AWS Service
Amazon VPC	Virtual Private Cloud enclosing the entire infrastructure.	Provides isolated network environment with subnets, security groups, and route tables for private/public access.	Amazon VPC
Amazon Route 53	DNS service handling domain resolution.	Routes client traffic to the nearest CloudFront edge or ELB endpoint; supports health checks for failover.	Amazon Route 53
Amazon CloudFront	Content Delivery Network (CDN).	Caches static/dynamic content at global edge locations to reduce latency and origin load.	Amazon CloudFront
Elastic Load Balancing (ELB)	Application Load Balancer (ALB) or Network Load Balancer (NLB).	Distributes incoming traffic across EC2 instances in the Apps cluster; supports SSL termination and sticky sessions.	Elastic Load Balancing

2. API and Ingress Layer

Component	Description	Purpose	AWS Service
API Gateway	Managed API proxy.	Exposes RESTful/HTTP APIs; handles request validation, transformation, authorization (e.g., JWT/OAuth), rate limiting, and integration with Lambda/EC2 backends.	Amazon API Gateway

3. Compute Layer (EC2 Apps)

Component	Description	Purpose	AWS Service
EC2 (Apps)	Cluster of EC2 instances running Kolate Predict Microservices.	Hosts stateless, scalable application logic for prediction and business workflows; auto-scales based on demand. Containerized (e.g., Docker/ECS).	Amazon EC2

4. Messaging and Orchestration Layer

Component	Description	Purpose	Technology/AWS Service

Kafka Core Microservices	Set of microservices acting as producers/consumers on Kafka topics.	Processes business events asynchronously; enables decoupling between services (e.g., predict -> provision).	Custom Microservices on EC2
Message Queues	Includes Zookeeper (for Kafka metadata/coordination) and Kafka brokers.	Reliable, partitioned event streaming; supports topics for core events like user actions, data syncs.	Apache Kafka + Zookeeper on EC2
Kolate Core	Central orchestration service for core microservices.	Coordinates workflows across Kafka consumers; likely includes saga patterns for distributed transactions.	Custom on EC2

5. Databases & Services Layer

Component	Description	Purpose	AWS Service/Technology
EC2 (Databases & Services)	Dedicated EC2 cluster for shared services and databases.	Runs stateful services like the Enterprise Info Provider; hosts core DBs with replication.	Amazon EC2
Enterprise Info Provider Service	Custom service for tenant onboarding/provisioning.	Manages enterprise data ingestion, transformation, and distribution to tenants; integrates with Kafka for events.	Custom Microservice on EC2
Kolate Core Databases	Shared datastores: - Redis DB : In-memory caching/key-value store. - Mongo DB : Document-oriented NoSQL for flexible schemas. - Postgres DB : Relational DB for structured data/transactions.	Core data persistence: Redis for sessions/caches, Mongo for logs/documents, Postgres for configs/users.	Redis, MongoDB, PostgreSQL on EC2 (or RDS equivalents)

Zookeeper	Distributed coordination service.	Manages Kafka cluster state, leader election, and configuration.	Apache Zookeeper on EC2
------------------	-----------------------------------	--	-------------------------

6. Tenant-Specific Layer

Component	Description	Purpose	AWS Service/Technology
EC2 (Tenant Databases)	Per-tenant EC2 clusters with containerized DBs.	Isolates tenant data for compliance (e.g., GDPR); scales independently per tenant.	Amazon EC2 with Containers (ECS/Fargate)
Tenant Databases	Replicated stack: - Postgres DB : Tenant transactions. - Mongo DB : Tenant documents. - Redis DB : Tenant caching.	Multi-tenant isolation; data sharded by tenant ID; supports sharding for large tenants.	PostgreSQL, MongoDB, Redis in Containers

Technologies and Integrations

- **Compute:** Primarily EC2 for flexibility; some components (e.g., tenant DBs) use containers for portability.
- **Messaging:** Apache Kafka (with Zookeeper) as the event backbone, ideal for high-throughput, real-time processing.
- **Databases:**
 - **Relational:** PostgreSQL for ACID compliance.
 - **NoSQL:** MongoDB for schema flexibility.
 - **Caching:** Redis for low-latency reads.
- **Networking:** Full AWS stack (VPC, Route 53, CloudFront, ELB, API Gateway) for global, secure delivery.
- **Microservices:** Likely built with Java/Node.js/Go; orchestrated via Kafka for event sourcing/CQRS patterns.
- **Monitoring/Scaling:** Implied but not shown—integrate with CloudWatch, Auto Scaling Groups.

Assumptions and Considerations

- **Scalability:** EC2 clusters should use Auto Scaling Groups tied to CPU/Memory metrics or Kafka lag.
- **Security:** Assume IAM roles for EC2, WAF on CloudFront/API Gateway, VPC endpoints for private traffic.

- **High Availability:** Multi-AZ deployment for EC2/Kafka/DBs; Route 53 failover routing.
- **Cost Optimization:** Use Reserved Instances for steady workloads; Spot Instances for non-critical microservices.
- **Deployment:** CI/CD via CodePipeline/CodeDeploy; Infrastructure as Code (IaC) with CloudFormation/Terraform.
- **Potential Enhancements:** Add AWS Lambda for serverless functions, S3 for blobs, or Elasticsearch for search.

This documentation provides a comprehensive view based on the diagram. For implementation details or code samples, additional context would be needed.